

PROGRAM 1:

AIM: To implement a program to load and view the dataset











ALGORITHM:

1. Download a dataset from Kaggle
2. Import the dataset or load in the jupyter book or in google colab
3. Create a duplicate set.
4. Display mean, median and other statistics using describe function().

PROGRAM:

```
import pandas as pd
import matplotlib.pyplot as plt
af = pd.read_csv("Most streamed spotify songs 2024.csv", encoding='unicode_escape')
new_af=af.copy()
new_af=new_af[["Artist", "release date", "All time rank", "spotify playlist reach", "youtube views"]]
new_af.head(10)
```

OUTPUT:

<div>2024-05-10</div> <div>@PopDataMusic</div> <div>Spotify Songs Global</div>									
#	↑/↓	Song	Credits	Streams		Peak		Days	Comment
				Amount	Change	#	Days		
1	NEW	 I Had Some Help (Feat. Morgan Wallen)	Post Malone, Morgan Wallen	13,949,573	NEW	1	1st	1	
2	↓-1	 Not Like Us	Kendrick Lamar	12,809,644	+3.47%	1	4	6	New Streams High
3	↓-1	 MILLION DOLLAR BABY	Tommy Richman	11,422,747	+3.33%	2	1	15	New Streams High
4	↓-1	 Espresso	Sabrina Carpenter	10,136,053	+3.73%	1	7	29	
5	↑5	 A Bar Song (Tipsy)	Shaboozey	7,604,741	+30.51%	5	NEW PEAK	26	New Streams High
6	↓-1	 Fortnight (feat. Post Malone)	Taylor Swift, Post Malone	7,440,880	+8.78%	1	10	22	
7	↓-3	 i like the way you kiss me	Artemas	7,284,624	+3.14%	1	14	52	
8	=	 Beautiful Things	Benson Boone	6,925,536	+8.7%	1	35	113	
9	↓-2	 Gata Only	FloyyMenor, Cris Mj	6,762,105	+4.43%	2	1	89	
10	↓-1	 Too Sweet	Hozier	6,747,550	+6.56%	2	10	50	

RESULT: Thus the desired output of most streamed songs has been successfully executed.

PROGRAM 2

AIM: To display the summary and statistics of the dataset.

ALGORITHM:

1. Download a dataset from Kaggle
2. Import the dataset or load in the jupyter book or in google colab
3. Create a duplicate set.
4. Display mean, median and other statistics using describe function().

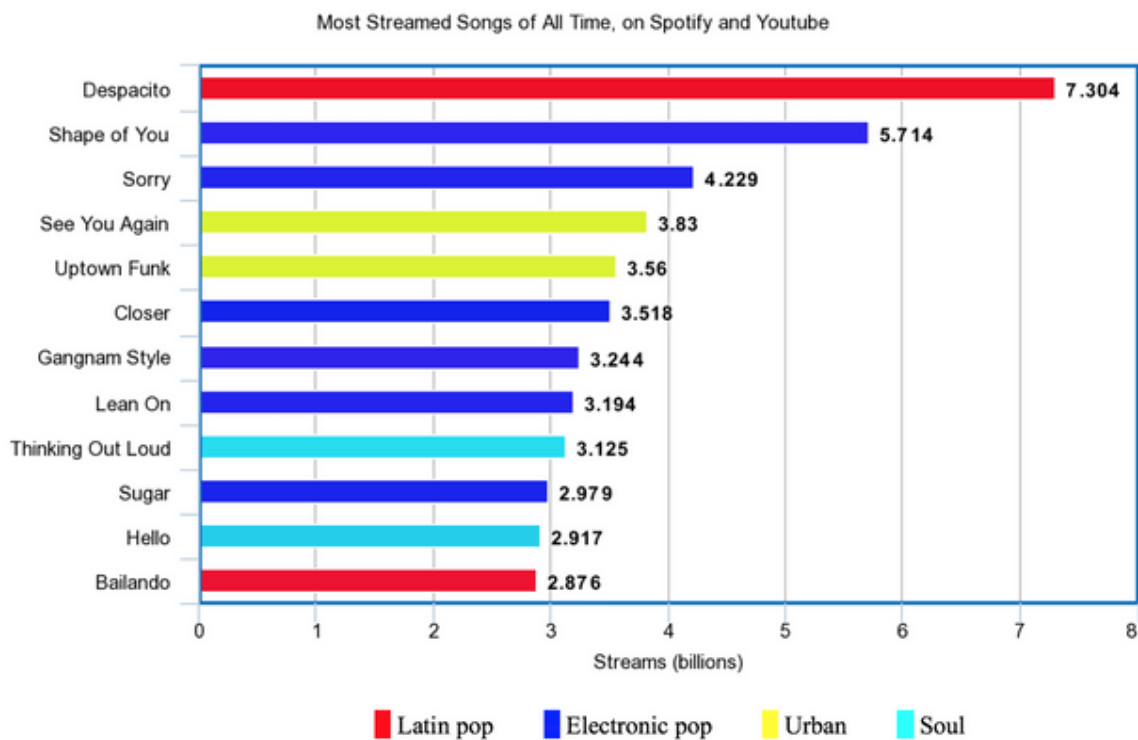
PROGRAM:

```
median = af["spotify popularity"].median()  
print("Median is ", median)
```

```
new_af.describe()
```

OUTPUT:

Median is 67.0



RESULT: Thus the program to display the summary and statistics has been successfully verified and executed.

PROGRAM 3

AIM: To implement linear regression to perform prediction.

ALGORITHM:

1. Initialize Parameters: Start by initializing the parameters like coefficients (slope and intercept).
2. Input Data: Gather the dataset containing the independent variable (X) and dependent variable (Y).
3. Feature Scaling (Optional): Normalize or standardize the input data if necessary to ensure better convergence.
4. Split Data: Divide the dataset into training and testing sets to evaluate the model.
5. Model Training: Implement a method to optimize the parameters (coefficients) based on the training data. This can be done using techniques like Gradient Descent, Normal Equations, or using libraries like scikit-learn.
6. Prediction: Use the learned parameters to predict outcomes for new data points or the test set.
7. Evaluation: Measure the performance of the model using evaluation metrics like Mean Squared Error (MSE), R-squared, etc.

PROGRAM:

```
import numpy as np
import matplotlib.pyplot as plt

# Sample dataset
X = np.array([1, 2, 3, 4, 5])
Y = np.array([2, 3, 4, 5, 6])

# Function to perform linear regression using Ordinary Least Squares
def linear_regression_ols(X, Y):

    # Add a column of ones to X for the intercept term
    X_b = np.c_[np.ones((len(X), 1)), X]

    # Calculate theta using the Normal Equation:  $\theta = (X^T * X)^{-1} * X^T * Y$ 
    theta = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(Y)
```

```

return theta

# Function to make predictions
def predict(X, theta):

    # Add a column of ones to X for the intercept term
    X_b = np.c_[np.ones((len(X), 1)), X]

    # Predict  $\hat{Y} = X_b * \theta$ 
    Y_pred = X_b.dot(theta)

    return Y_pred

# Perform linear regression
theta = linear_regression_ols(X, Y)

# Make predictions
X_new = np.array([6, 7])
predictions = predict(X_new, theta)

# Plotting the original data and the linear regression line
plt.scatter(X, Y, color='blue', label='Data points')

plt.plot(X_new, predictions, color='red', label='Linear Regression')
plt.xlabel('X')


plt.ylabel('Y')
plt.legend()

plt.title('Linear Regression using Ordinary Least Squares')
plt.show()

print(f"Predictions for X_new: {predictions}")

```

OUTPUT:



RESULT: Thus the implementation of linear regression to perform prediction has been successfully executed.

PROGRAM 4.1

AIM: To implement Bayesian logistic regression for classification

ALGORITHM:

8. **Initialize Parameters:** Start with prior distributions for the model parameters (e.g., coefficients, intercept) and likelihood distributions based on the data.
9. **Input Data:** Gather the dataset containing features (X) and corresponding binary labels (Y).
10. **Posterior Calculation:** Use Bayesian inference techniques such as Markov Chain Monte Carlo (MCMC) or Variational Inference to compute the posterior distribution over the parameters given the data.
11. **Prediction:** Use the posterior distribution to predict the probability of classes for new data points.
12. **Evaluation:** Measure the performance of the model using metrics such as accuracy, precision, recall, and F1-score.

PROGRAM:

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Generate synthetic data
X, Y = make_classification(n_samples=1000, n_features=20, random_state=42)

# Split data into train and test sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

# Add intercept to X_train for bias term
X_train = np.c_[np.ones((len(X_train), 1)), X_train]
```

```

# Define sigmoid function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Initialize parameters randomly
np.random.seed(42)
theta = np.random.randn(X_train.shape[1])

# Bayesian Logistic Regression with Metropolis-Hastings sampling
def bayesian_logistic_regression(X, Y, num_samples=1000, burn_in=200):
    m, n = X.shape
    trace = np.zeros((num_samples, n)) # Trace to store samples of theta
    theta_current = theta.copy()
    acceptance_count = 0

    for i in range(num_samples):
        # Generate proposal from Gaussian distribution
        proposal = theta_current + np.random.randn(n)

        # Calculate prior probabilities (assuming uninformative priors)
        prior_current = np.sum(-np.log(1 + np.exp(-(X.dot(theta_current)))))
        prior_proposal = np.sum(-np.log(1 + np.exp(-(X.dot(proposal)))))

        # Calculate likelihoods
        likelihood_current = np.sum(Y * X.dot(theta_current) - np.log(1 +
np.exp(X.dot(theta_current))))
        likelihood_proposal = np.sum(Y * X.dot(proposal) - np.log(1 + np.exp(X.dot(proposal))))

        # Calculate posterior probabilities
        posterior_current = likelihood_current + prior_current

```

```
posterior_proposal = likelihood_proposal + prior_proposal
```

```
# Accept or reject the proposal
```

```
acceptance_prob = np.exp(posterior_proposal - posterior_current)
```

```
accept = np.random.rand() < acceptance_prob
```

```
if accept:
```

```
    theta_current = proposal
```

```
    acceptance_count += 1
```

```
trace[i] = theta_current
```

```
acceptance_rate = acceptance_count / num_samples
```

```
print(f'Acceptance rate: {acceptance_rate}')
```

```
return trace[burn_in:]
```

```
# Perform Bayesian Logistic Regression
```

```
trace = bayesian_logistic_regression(X_train, Y_train)
```

```
# Predictions for test data
```

```
X_test = np.c_[np.ones((len(X_test), 1)), X_test]
```

```
logits = X_test.dot(trace.mean(axis=0))
```

```
Y_pred = (sigmoid(logits) >= 0.5).astype(int)
```

```
# Calculate accuracy
```

```
accuracy = accuracy_score(Y_test, Y_pred)
```

```
print(f'Accuracy: {accuracy}')
```

```
# Plotting the coefficients distribution
```

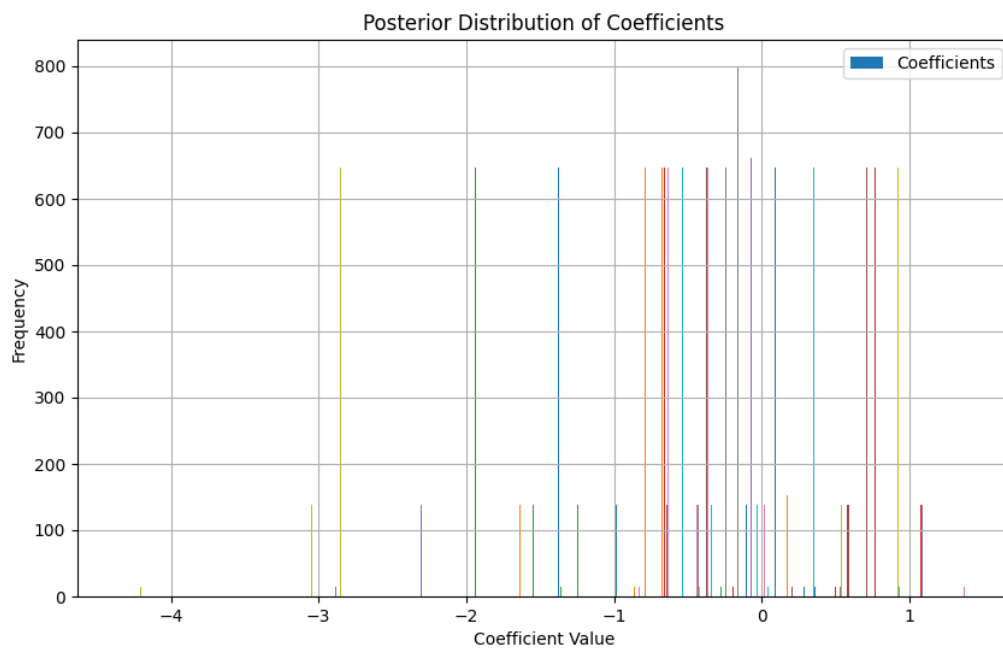


```
plt.figure(figsize=(10, 6))
plt.hist(trace[:, 1:], bins=30, label='Coefficients')
plt.xlabel('Coefficient Value')
plt.ylabel('Frequency')
plt.title('Posterior Distribution of Coefficients')
plt.legend()
plt.grid(True)
plt.show()
```

OUTPUT:

Acceptance rate: 0.005

Accuracy: 0.69



RESULT: Thus the program for to implement Bayesian logistic regression for classification is been successfully executed.

PROGRAM 4.2

AIM: To implement the SVM for classification

ALGORITHM:

- i. **Initialize Parameters:** Start with setting parameters such as kernel type (linear, polynomial, radial basis function (RBF)), regularization parameter (C), and kernel coefficients (gamma).
- ii. **Input Data:** Gather the dataset containing features (X) and corresponding binary labels (Y).
- iii. **Model Training:** Use the training data to fit the SVM model, adjusting the parameters to maximize the margin between classes.
- iv. **Prediction:** Use the learned SVM model to predict classes for new data points.
- v. **Evaluation:** Measure the performance of the model using metrics such as accuracy, precision, recall, and F1-score.

PROGRAM:

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.datasets import make_classification

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, classification_report

# Generate synthetic data

X, Y = make_classification(n_samples=1000, n_features=20, random_state=42)

# Split data into train and test sets

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

# Define SVM model

svm_model = SVC(kernel='rbf', C=1.0, gamma='scale', random_state=42)

# Train SVM model
```

```

svm_model.fit(X_train, Y_train)

# Predictions for test data
Y_pred = svm_model.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(Y_test, Y_pred)
print(f'Accuracy: {accuracy}')

# Classification report
print(classification_report(Y_test, Y_pred))

# Plotting decision boundary (for 2D data)
if X.shape[1] == 2:
    # Plot decision boundary
    plt.figure(figsize=(8, 6))
    plt.scatter(X[:, 0], X[:, 1], c=Y, cmap='viridis', s=50, alpha=0.6)

    # Create grid to evaluate model
    xlim = plt.gca().get_xlim()
    ylim = plt.gca().get_ylim()
    xx, yy = np.meshgrid(np.linspace(xlim[0], xlim[1], 100),
                          np.linspace(ylim[0], ylim[1], 100))
    Z = svm_model.decision_function(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot decision boundary and margins
    plt.contour(xx, yy, Z, colors='k', levels=[-1, 0, 1], alpha=0.5,
                linestyle=['--', '-', '--'])

```

```
plt.title('SVM Decision Boundary')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```

OUTPUT:

Accuracy: 0.845

	precision	recall	f1-score	support
0	0.80	0.89	0.84	93
1	0.90	0.80	0.85	107
accuracy		0.84		200
macro avg	0.85	0.85	0.84	200
weighted avg	0.85	0.84	0.85	200

RESULT: The given program for SVM classification is executed and verified successfully.

PROGRAM 5.1

AIM : To implement the K-means clustering to categorize the data

ALGORITHM:

1. **Initialize:** Randomly select KKK initial centroids from the data.
2. **Assignment:** Assign each data point to the nearest centroid, forming KKK clusters.
3. **Update:** Recalculate the centroids of the clusters by taking the mean of all data points in each cluster.
4. **Repeat:** Repeat steps 2 and 3 until the centroids no longer change (convergence) or for a fixed number of iterations.

PROGRAM:

```
import numpy as np

import matplotlib.pyplot as plt

# Generate synthetic data

def generate_data(n_samples=300, n_centers=4, random_seed=42):

    np.random.seed(random_seed)

    points_per_center = n_samples // n_centers

    centers = np.random.uniform(-10, 10, (n_centers, 2))

    X = np.vstack([center + np.random.randn(points_per_center, 2) for center in centers])

    return X

# K-Means Clustering

def k_means(X, k, max_iters=100):

    centroids = X[np.random.choice(X.shape[0], k, replace=False)]

    for _ in range(max_iters):

        # Assign each point to the nearest centroid

        distances = np.linalg.norm(X[:, np.newaxis] - centroids, axis=2)
```

```
clusters = np.argmin(distances, axis=1)

# Calculate new centroids
new_centroids = np.array([X[clusters == j].mean(axis=0) for j in range(k)])

# Check for convergence
if np.all(centroids == new_centroids):
    break

centroids = new_centroids

return clusters, centroids

# Generate data and run K-Means
X = generate_data()
clusters_kmeans, centroids_kmeans = k_means(X, k=4)

# Plot K-Means results
plt.figure(figsize=(10, 6))


plt.scatter(X[:, 0], X[:, 1], c=clusters_kmeans, cmap='viridis', marker='o')
plt.scatter(centroids_kmeans[:, 0], centroids_kmeans[:, 1], c='red', marker='x')

plt.title('K-Means Clustering')

plt.xlabel('Feature 1')
plt.ylabel('Feature 2')

plt.show()
```

OUTPUT:



RESULT: Thus the program for K-means clustering is been executed successfully

PROGRAM 5.2

AIM: To implement the mixture of gaussian models to categorize the data

ALGORITHM:

1. **Initialize:** Choose initial parameters for the Gaussian components (means, covariances, and mixing coefficients).
2. **Expectation (E-step):** Calculate the probability of each data point belonging to each Gaussian component.
3. **Maximization (M-step):** Update the parameters of the Gaussian components using the probabilities computed in the E-step.
4. **Repeat:** Repeat the E-step and M-step until convergence.

PROGRAM:

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
# Generate synthetic data
```

```
def generate_data(n_samples=300, n_centers=4, random_seed=42):
```

```
    np.random.seed(random_seed)
```

```
    points_per_center = n_samples // n_centers
```

```
    centers = np.random.uniform(-10, 10, (n_centers, 2))
```

```
    X = np.vstack([center + np.random.randn(points_per_center, 2) for center in centers])
```

```
    return X
```

```
# Gaussian Mixture Model (GMM)
```

```
def gmm(X, k, max_iters=100):
```

```
    n_samples, n_features = X.shape
```

```
    # Initialize parameters
```

```
    np.random.seed(42)
```

```
    weights = np.ones(k) / k
```

```
    means = X[np.random.choice(X.shape[0], k, replace=False)]
```

```
    covariances = np.array([np.eye(n_features) for _ in range(k)])
```

```
def gaussian(x, mean, cov):
```

```
    n = x.shape[0]
```

```
    diff = x - mean
```

```
    return (np.exp(-0.5 * np.dot(diff.T, np.linalg.solve(cov, diff))) /
```

```
        np.sqrt((2 * np.pi) ** n * np.linalg.det(cov)))
```

```
def e_step(X, weights, means, covariances):
```



```

responsibilities = np.zeros((n_samples, k))

for i in range(n_samples):

    for j in range(k):

        responsibilities[i, j] = weights[j] * gaussian(X[i], means[j], covariances[j])

    responsibilities[i] /= np.sum(responsibilities[i])

return responsibilities

```

```

def m_step(X, responsibilities):

    weights = np.mean(responsibilities, axis=0)

    means = np.dot(responsibilities.T, X) / np.sum(responsibilities, axis=0)[:, np.newaxis]

    covariances = []

    for j in range(k):

        diff = X - means[j]

        cov = np.dot(responsibilities[:, j] * diff.T, diff) / np.sum(responsibilities[:, j])

        covariances.append(cov)

    return weights, means, np.array(covariances)

```

```

for _ in range(max_iters):

    responsibilities = e_step(X, weights, means, covariances)

    weights, means, covariances = m_step(X, responsibilities)

clusters = np.argmax(responsibilities, axis=1)

return clusters, means

```

Generate data and run GMM

```
X = generate_data()

clusters_gmm, means_gmm = gmm(X, k=4)

# Plot GMM results

plt.figure(figsize=(10, 6))

plt.scatter(X[:, 0], X[:, 1], c=clusters_gmm, cmap='viridis', marker='o')

plt.scatter(means_gmm[:, 0], means_gmm[:, 1], c='red', marker='x')


plt.title('Gaussian Mixture Model Clustering')

plt.xlabel('Feature 1')

plt.ylabel('Feature 2')

plt.show()
```

OUTPUT:



RESULT:

Thus the program is been executed successfully.

PROGRAM 5.3

AIM: To implement the hierarchial clustering to categorize the data

ALGORITHM:

1. **Start:** Treat each data point as a singleton cluster.
2. **Merge:** Find the pair of clusters that are closest and merge them into a single cluster.

3. **Repeat:** Repeat step 2 until only a single cluster remains or a stopping criterion is met (e.g., a desired number of clusters).
4. **Cut:** Cut the dendrogram at the desired level to extract clusters.

PROGRAM:

```
import numpy as np

import matplotlib.pyplot as plt


# Generate synthetic data

def generate_data(n_samples=300, n_centers=4, random_seed=42):

    np.random.seed(random_seed)

    points_per_center = n_samples // n_centers

    centers = np.random.uniform(-10, 10, (n_centers, 2))

    X = np.vstack([center + np.random.randn(points_per_center, 2) for center in centers])

    return X


# Calculate Euclidean distance

def euclidean_distance(a, b):

    return np.sqrt(np.sum((a - b) ** 2))


# Compute the distance matrix

def compute_distance_matrix(X):

    n_samples = X.shape[0]

    distances = np.zeros((n_samples, n_samples))

    for i in range(n_samples):

        for j in range(i + 1, n_samples):
```

```
distances[i, j] = euclidean_distance(X[i], X[j])

distances[j, i] = distances[i, j]

return distances
```

```
# Hierarchical Clustering
```

```
def hierarchical_clustering(X):
```

```
    distances = compute_distance_matrix(X)
```

```
    n_samples = len(X)
```

```
    # Initialize clusters
```

```
    clusters = [[i] for i in range(n_samples)]
```

```
    while len(clusters) > 1:
```

```
        # Find the two closest clusters
```

```
        min_dist = float('inf')
```

```
        to_merge = (None, None)
```

```
        for i in range(len(clusters)):
```

```
            for j in range(i + 1, len(clusters)):
```

```
                d = np.min([distances[p][q] for p in clusters[i] for q in clusters[j]])
```

```
                if d < min_dist:
```

```
                    min_dist = d
```

```
                    to_merge = (i, j)
```

```
        # Merge the two clusters
```

```

        i, j = to_merge

        clusters[i].extend(clusters[j])

        del clusters[j]

    return clusters

# Function to extract cluster labels
def extract_clusters(clusters, n_samples):

    labels = np.zeros(n_samples)

    for cluster_id, cluster in enumerate(clusters):

        for index in cluster:

            labels[index] = cluster_id

    return labels

# Generate data and perform hierarchical clustering
X = generate_data()

final_clusters = hierarchical_clustering(X)

# Extract final cluster labels
cluster_labels = extract_clusters(final_clusters, len(X))

# Plot Hierarchical Clustering results
plt.figure(figsize=(10, 6))

plt.scatter(X[:, 0], X[:, 1], c=cluster_labels, cmap='viridis', marker='o')


plt.title('Hierarchical Clustering')
```

```
plt.xlabel('Feature 1')
```

```
plt.ylabel('Feature 2')
```

```
plt.show()
```

OUTPUT:



RESULT: Thus the desired program have been successfully executed .

PROGRAM 6

AIM: To create a program to perform PCA

ALGORITHM:

1. **Standardize the Data:** Center the data by subtracting the mean of each feature from the data. Optionally, scale each feature to unit variance.
2. **Compute the Covariance Matrix:** Calculate the covariance matrix of the centered data.
3. **Calculate Eigenvalues and Eigenvectors:** Find the eigenvalues and eigenvectors of the covariance matrix. The eigenvectors determine the directions of the new feature space, and the eigenvalues determine their magnitude.
4. **Sort Eigenvalues and Eigenvectors:** Sort the eigenvectors by decreasing eigenvalues and select the top k eigenvectors to form a matrix (principal components).
5. **Transform the Data:** Project the original data onto the new feature space using the matrix of principal components.

PROGRAM:

```
import numpy as np

def pca(X, n_components):

    # Step 1: Center the Data

    X_centered = X - np.mean(X, axis=0)

    # Step 2: Compute the Covariance Matrix

    cov_matrix = np.cov(X_centered, rowvar=False)

    # Step 3: Calculate Eigenvalues and Eigenvectors

    eigenvalues, eigenvectors = np.linalg.eigh(cov_matrix)

    # Step 4: Sort Eigenvalues and Eigenvectors

    sorted_index = np.argsort(eigenvalues)[::-1]
```



```
sorted_eigenvectors = eigenvectors[:, sorted_index]
```

```
sorted_eigenvalues = eigenvalues[sorted_index]
```

```
# Step 5: Select Top n_components Eigenvectors
```

```
eigenvector_subset = sorted_eigenvectors[:, :n_components]
```

```
# Step 6: Transform the Data
```

```
X_reduced = np.dot(X_centered, eigenvector_subset)
```

```
return X_reduced, sorted_eigenvalues, eigenvector_subset
```

```
# Example usage
```

```
if __name__ == "__main__":
```

```
    # Example data
```

```
    X = np.array([[2.5, 2.4],
```

```
                  [0.5, 0.7],
```

```
                  [2.2, 2.9],
```

```
                  [1.9, 2.2],
```

```
                  [3.1, 3.0],
```

```
                  [2.3, 2.7],
```

```
                  [2, 1.6],
```

```
                  [1, 1.1],
```

```
                  [1.5, 1.6],
```

```
                  [1.1, 0.9]])
```

```
# Perform PCA

X_reduced, eigenvalues, eigenvectors = pca(X, n_components=2)

print("Reduced Data:\n", X_reduced)

print("Eigenvalues:\n", eigenvalues)

print("Eigenvectors:\n", eigenvectors)
```

OUTPUT:

Reduced Data:

```
[[ 0.82797019 -0.17511531]
 [-1.77758033  0.14285723]
 [ 0.99219749  0.38437499]
 [ 0.27421042  0.13041721]
 [ 1.67580142 -0.20949846]
 [ 0.9129491  0.17528244]
 [-0.09910944 -0.3498247 ]
 [-1.14457216  0.04641726]
 [-0.43804614  0.01776463]
 [-1.22382056 -0.16267529]]
```

Eigenvalues:

```
[1.28402771 0.0490834 ]
```

Eigenvectors:

```
[[ 0.6778734 -0.73517866]
 [ 0.73517866  0.6778734 ]]
```

RESULT: Thus the given program has been successfully executed

PROGRAM 7

AIM : To implement HMM to predict the sequential data.

ALGORITHM :

```
import numpy as np
```

```
class SimpleHMM:
```

```
    def __init__(self, states, observations, start_prob, trans_prob, emit_prob):
```

```
        self.states = states
```

```
        self.observations = observations
```

```
        self.start_prob = start_prob
```

```
        self.trans_prob = trans_prob
```

```
        self.emit_prob = emit_prob
```

```
    def viterbi(self, obs_sequence):
```

```
        n_states = len(self.states)
```

```
        n_obs = len(obs_sequence)
```

```
        # Initialize the dynamic programming tables
```

```
        viterbi_table = np.zeros((n_states, n_obs))
```

```
        backpointer_table = np.zeros((n_states, n_obs), dtype=int)
```

```
        # Initialization step
```

```
        first_obs = obs_sequence[0]
```

```
        for s in range(n_states):
```

```
            viterbi_table[s, 0] = self.start_prob[s] * self.emit_prob[s, first_obs]
```

```

    backpointer_table[s, 0] = 0

# Recursion step

for t in range(1, n_obs):

    for s in range(n_states):

        probabilities = viterbi_table[:, t-1] * self.trans_prob[:, s] * self.emit_prob[s,
obs_sequence[t]]

        viterbi_table[s, t] = np.max(probabilities)

        backpointer_table[s, t] = np.argmax(probabilities)

# Termination step

best_path_prob = np.max(viterbi_table[:, n_obs-1])

best_last_state = np.argmax(viterbi_table[:, n_obs-1])

# Path backtracking

best_path = np.zeros(n_obs, dtype=int)

best_path[-1] = best_last_state

for t in range(n_obs-2, -1, -1):

    best_path[t] = backpointer_table[best_path[t+1], t+1]

return best_path, best_path_prob

# Example usage

if __name__ == "__main__":

    # Define the states, observations, and model parameters

```

```

states = ['Rainy', 'Sunny']

observations = ['Walk', 'Shop', 'Clean']

start_probability = np.array([0.6, 0.4])

transition_probability = np.array([
    [0.7, 0.3], # From Rainy to Rainy/Sunny
    [0.4, 0.6], # From Sunny to Rainy/Sunny
])

emission_probability = np.array([
    [0.1, 0.4, 0.5], # Probabilities of Walk/Shop/Clean from Rainy
    [0.6, 0.3, 0.1], # Probabilities of Walk/Shop/Clean from Sunny
])

# Create the HMM model

hmm = SimpleHMM(states, observations, start_probability, transition_probability,
emission_probability)

# Encode the observation sequence as integers

obs_map = {obs: i for i, obs in enumerate(observations)}

obs_sequence = np.array([obs_map['Walk'], obs_map['Shop'], obs_map['Clean'],
obs_map['Walk']])

# Predict the most likely sequence of states

```

```
state_sequence, probability = hmm.viterbi(obs_sequence)
```

```
# Decode state sequence into state names
```

```
state_names = [states[state] for state in state_sequence]
```

```
print("Most likely states sequence:", state_names)
```

```
print("Probability of the best path:", probability)
```

OUTPUT:

Most likely states sequence: ['Sunny', 'Rainy', 'Rainy', 'Sunny']

Probability of the best path: 0.0024192

RESULT: Thus the program to implement HMM has been successfully executed .

PROGRAM 8

AIM: To implement the CART learning algorithms to perform categorization.

ALGORITHM:

```
from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn import tree

import matplotlib.pyplot as plt


# Load the Iris dataset

data = load_iris()

X, y = data.data, data.target


# Split the dataset into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)


# Create a Decision Tree classifier

clf = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=42)


# Train the classifier

clf.fit(X_train, y_train)
```

```
# Make predictions on the test set
```

```
y_pred = clf.predict(X_test)
```

```
# Evaluate the model
```

```
accuracy = clf.score(X_test, y_test)
```

```
print(f'Accuracy: {accuracy:.2f}')
```

```
# Visualize the Decision Tree
```

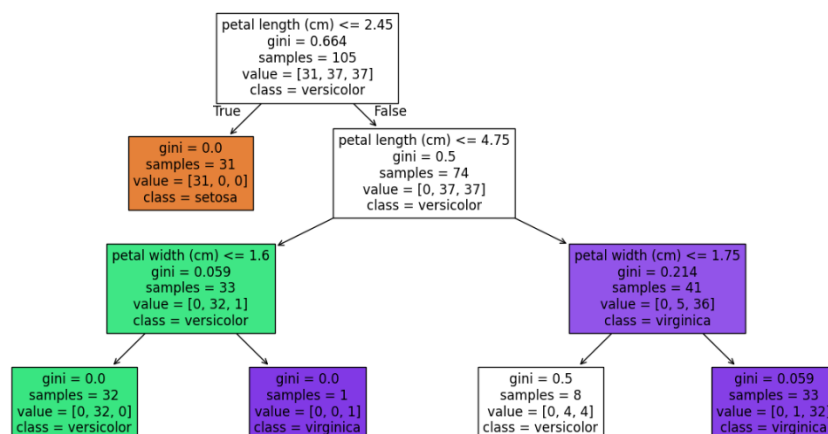
```
plt.figure(figsize=(12, 8))
```

```
tree.plot_tree(clf, feature_names=data.feature_names, class_names=data.target_names, filled=True)
```

```
plt.show()
```

OUTPUT:

Accuracy:1.0



RESULT: Thus the given program for the implementation of CART learning algorithms to perform categorization has been successfully executed.

PROGRAM 9

AIM: To implement Ensemble learning models to perform classification

ALGORITHM:

Bagging (e.g., Random Forest)

1. **Create multiple subsets of the training data** by sampling with replacement (bootstrap sampling).
2. **Train a base model** (e.g., decision tree) on each subset independently.
3. **Aggregate the predictions** of all models by taking a majority vote (for classification) or averaging (for regression).

Boosting (e.g., AdaBoost)

1. **Initialize weights** for all training examples.
2. **Train a base model** on the training data, weighted according to their current weights.
3. **Evaluate the model** and increase the weights of misclassified examples.
4. **Train subsequent models** iteratively, focusing more on difficult examples.
5. **Combine the models** by giving more weight to better-performing models.

PROGRAM:

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier

from sklearn.decomposition import PCA
```

```
# Load Iris dataset and split into train and test sets

X, y = load_iris(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)


# Reduce dimensionality for visualization purposes

pca = PCA(n_components=2)

X_train_2d = pca.fit_transform(X_train)

X_test_2d = pca.transform(X_test)


# Train Random Forest classifier

rf_clf = RandomForestClassifier(random_state=42)

rf_clf.fit(X_train_2d, y_train)


# Train AdaBoost classifier

ada_clf = AdaBoostClassifier(random_state=42)

ada_clf.fit(X_train_2d, y_train)


# Function to plot decision boundaries

def plot_decision_boundaries(clf, X, y, ax, title):

    h = .02 # Step size in the mesh

    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1

    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),

                          np.arange(y_min, y_max, h))
```

```
Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])  
  
Z = Z.reshape(xx.shape)  
  
ax.contourf(xx, yy, Z, alpha=0.3)  
  
ax.scatter(X[:, 0], X[:, 1], c=y, edgecolor='k', marker='o')  
  
ax.set_title(title)
```

```
# Create plots
```

```
fig, axs = plt.subplots(1, 2, figsize=(14, 6))
```

```
# Plot Random Forest decision boundaries
```


```
plot_decision_boundaries(rf_clf, X_test_2d, y_test, axs[0], 'Random Forest')
```

```
# Plot AdaBoost decision boundaries
```

```
plot_decision_boundaries(ada_clf, X_test_2d, y_test, axs[1], 'AdaBoost')
```

```
plt.show()
```

OUTPUTS:



RESULT: Thus the program for implementation to Ensemble learning models to perform classification is been successfully executed and verified.